Supporting Information for: Hard Shape Constraints versus Penalties for Temperature Extrapolation of MOF and COF Water Isotherms

Abdulhamid Batayhi
Department of Mechatronics Engineering, Faculty of Mechanical Engineering,
Yildiz Technical University, 34220 Istanbul, Türkiye
E-mail: abdulhamid.batayhi@std.yildiz.edu.tr

Contents: Tables S1–S8; Figure S1; Notes S1–S4 (synthetic worlds and power; synthetic truths for the isosteric heat; digitisation error; model definitions); references.

Table S1: Data

Material	Class	Set	Data	Source DOI	T range (K)	n_T	Points per T	p/p_{sat} range	q_{ref} (mol kg^{-1})
COF-432	COF	conf.	vector	10.1021/jacs.9b13094	283.15–313.15	3	28–29	0.063–0.870	16.04
COF-482-hydrazide	COF	conf.	raster	10.1021/acscentsci.2c00398	288.15–308.15	3	25–51	0.08–0.898	34.70
COF-SO_3H	COF	conf.	vector	10.1002/cssc.202301906	298.15–318.15	3	18–18	0.0003–0.900	17.04
HFPTP-DMePDA-COF	COF	conf.	raster	10.1038/s41467-021-27128-4	283.15–298.15	3	9–18	0.39–0.897	7.91
HFPTP-PDA-COF	COF	conf.	raster	10.1038/s41467-021-27128-4	283.15–298.15	3	24–32	0.0041–0.898	16.78
NO-PI-3-COF	COF	conf.	vector	10.1021/jacs.3c02572	288.15–308.15	3	30–37	0.0091–0.882	12.74
PI-3-COF	COF	conf.	vector	10.1021/jacs.3c02572	288.15–308.15	3	28–37	0.01–0.887	24.80
TFBCz-PDA-COF	COF	conf.	raster	10.1038/s41467-021-27128-4	283.15–298.15	3	29–30	0.0023–0.899	27.26
Al-fumarate	MOF	conf.	raster	10.1021/acscentsci.9b00745	283.15–323.15	3	56–59	0.049–0.892	23.86
CALF-25	MOF	conf.	ISODB	10.1021/ja306812r	298.00–313.00	3	17–20	0.018–0.894	3.75
Co_2Cl_2(BTDD)	MOF	conf.	vector	10.1021/acscentsci.7b00186	283.00–298.00	3	31–34	0.0072–0.862	53.97
Mg-SU-102	MOF	conf.	numeric	10.1038/s41467-025-59551-2	288.15–328.15	9	21–21	0.0013–0.819	21.11
Zn_2(bptc)	MOF	conf.	ISODB	10.1021/ja060946u	288.00–318.00	4	25–33	0.023–0.885	8.38
TFPPy-PDA-COF	COF	dev.	raster	10.1038/s41467-021-27128-4	283.15–298.15	3	7–15	0.46–0.900	35.41
CuBTC	MOF	dev.	ISODB	10.1016/j.ijheatmasstransfer.2012.07.068	293.00–358.00	9	23–41	0.0029–0.508	25.29
MIL-100	MOF	dev.	ISODB	10.1016/j.ijheatmasstransfer.2012.07.068	293.00–358.00	9	25–35	0.011–0.603	18.11
MIL-160	MOF	dev.	ISODB	10.1002/adma.201502418	293.00–313.00	3	32–54	0.0052–0.884	21.79

Table S1. The 13 confirmatory materials (conf.) and the four development materials (dev.) after the inclusion rules and the preprocessing. Data: how the points were obtained (vector or raster digitisation of the published figure, numeric source data, or the NIST adsorption database, ISODB). n_T: number of temperatures; points per temperature with $p/p_{\mathrm{sat}} \leq 0.90$; q_{ref}: the largest uptake over all temperatures in this range, used only for scoring. Every decision and its reason is in data/inclusion_log.csv.

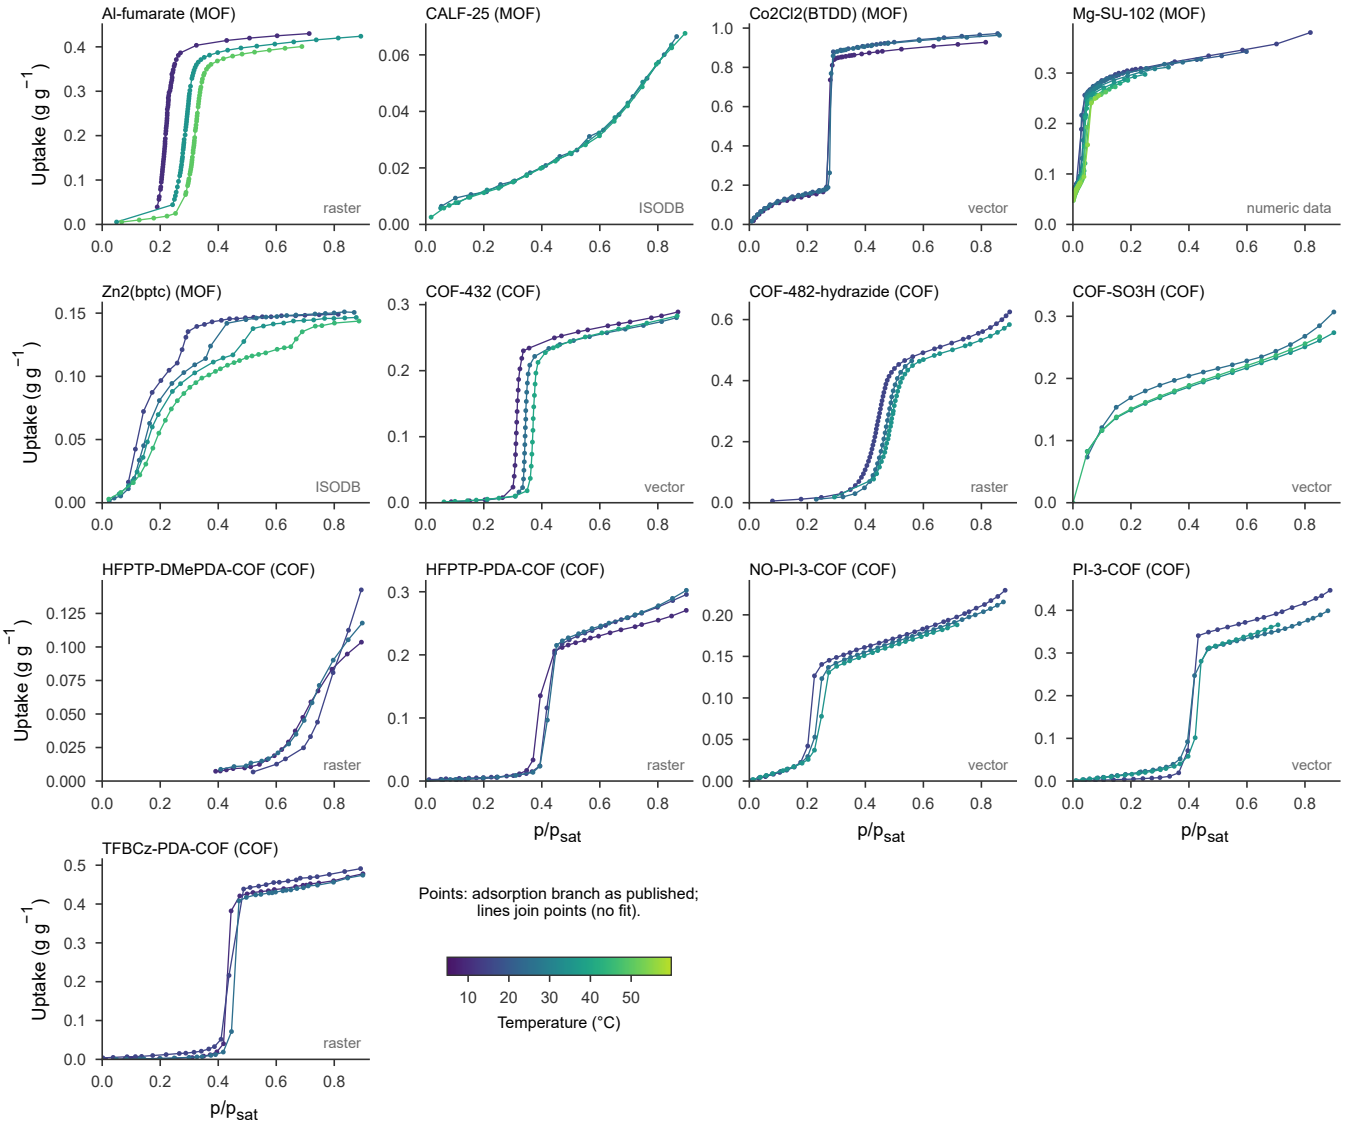


Figure S1. The confirmatory data: the adsorption isotherms of the 13 materials at every temperature, as used (relative pressure up to 0.90), coloured by temperature; points as published, joined by thin lines with no fit. The label in each panel gives how the points were obtained.

Tables S2 to S7: Results in detail

			$R_m = E_P/E_X$								
Material	Class	n_{train}	H1b	V-pen	H1c	A-hard	H1d	F-none	H1e	C1	H1a	F-hard†	V-none†	F-rel†
COF-432	COF	2	0.86	0.83	0.37	0.98	0.70	0.84	1.58					
COF-482-hydrazide	COF	2	0.95	1.18	0.35	1.03	0.22	0.96	0.72					
COF-SO_3H	COF	2	10.91	9.92	6.19	8.58	6.04	9.80	5.35					
HFPTP-DMePDA-COF	COF	2	1.23	4.60	0.51	0.74	1.14	0.94	0.95					
HFPTP-PDA-COF	COF	2	0.99	2.28	1.01	1.01	0.85	1.00	0.95					
NO-PI-3-COF	COF	2	2.47	1.55	1.07	2.47	0.20	2.09	2.69					
PI-3-COF	COF	2	1.10	4.81	1.50	2.20	0.22	1.45	1.99					
TFBCz-PDA-COF	COF	2	0.78	3.14	0.99	0.97	0.39	0.94	0.83					
Al-fumarate	MOF	2	0.87	0.49	0.28	0.90	0.39	0.94	1.19					
CALF-25	MOF	2	0.94	0.98	0.39	0.59	0.77	0.90	0.41					
Co_2Cl_2(BTDD)	MOF	2	1.44	1.13	1.02	0.92	0.59	1.38	1.92					
Mg-SU-102	MOF	8	1.01	0.73	0.24	1.56	0.28	1.07	0.31					
Zn_2(bptc)	MOF	3	1.27	0.32	0.66	1.66	0.49	1.34	0.67					
Hodges-Lehmann R			1.10	1.51	0.63	1.23	0.49	1.14	1.13					

Table S2. Error ratio of P against each comparator per material, $R_m = 10^{d_m}$, with d_m the difference of the mean of $\log_{10} E$ over E-high and E-low; below one, P extrapolated better. n_{train}: the smallest number of training temperatures over the two tasks. The last row is the Hodges-Lehmann estimate over the 13 materials (main text, Figure 2a). Columns marked † (H1a, V-none, F-rel) are descriptive comparisons.

		Shape violations (H3)						Isosteric heat (fits with any point)			
Model	Fits	In domain	Wilson 95 %	Band	Strict	All tasks	Class	Undefined	Outside 0–150	< 0	> 150
P	26	0	0.00–0.13	0	0	0/39	none	0	14	10	6
F-hard	26	0	0.00–0.13	0	0	0/39	none	0	23	20	16
A-hard	26	0	0.00–0.13	0	0	0/39	none	0	0	0	0
C1	26	0	0.00–0.13	0	0	0/39	none	0	2	1	1
P-K	2	0	0.00–0.66	0	0	0/3	none	0	2	–	–
V-pen	26	9	0.19–0.54	10	11	12/39	some	11	12	12	5
V-none	26	8	0.17–0.50	9	10	12/39	some	8	15	11	8
F-pen	26	13	0.32–0.68	12	24	19/39	often	23	26	–	–
A-pen	26	16	0.43–0.78	16	21	25/39	often	16	0	–	–
A-none	26	15	0.39–0.74	15	19	24/39	often	17	0	–	–
F-rel	26	22	0.66–0.94	22	26	34/39	often	25	26	–	–
F-none	26	26	0.87–1.00	26	26	39/39	often	25	26	24	26

Table S3. Violating fits per model over the 26 fits of E-high and E-low. Shape violations: In domain, the pre-registered H3 count, a fall of more than 2 % of q_{ref} below the running maximum or an uptake below -2 % of q_{ref}, inside the penalty domain; Wilson 95 % interval of its rate; Band, the same rule on the rest of the grid, up to 20 K beyond the measured temperatures and to 0.95 p_{sat}; Strict, any decrease of uptake with pressure or any negative uptake anywhere on the grid, with no tolerance; All tasks, the in-domain count adding I-mid; Class, the pre-registered word. Isosteric heat (pre-registered strict counts, omitted from the first draft; ledger X-5): fits with any grid point where the heat is undefined, or outside 0 to 150 $\mathrm{kJ\,mol^{-1}}$, from the run records; that count split into negative heats and heats above 150 $\mathrm{kJ\,mol^{-1}}$ from exact refits of the reported seeds (ledger X-4; computed for the models in the H1 comparisons and V-none, – otherwise).

Model	Fits	Median range	Largest range	Range above 1	Selected best	Selected worst
P	26	0.04	2.34	1	6	9
V-pen	26	0.11	1.46	1	5	5
V-none	26	0.13	1.76	2	10	7
A-hard	26	0.01	0.28	0	12	5
A-pen	26	0.03	0.54	0	7	4
A-none	26	0.04	1.41	2	6	4
C1	26	0.01	2.84	1	4	2
F-hard	26	0.54	1.48	5	7	7
F-pen	26	0.67	1.28	3	4	2
F-none	26	0.26	0.95	0	8	5
F-rel	26	0.18	0.97	0	3	10
P-K	2	0.34	0.47	0	0	1

Table S4. Seed spread of the held-out error over the E-high and E-low fits. Range: the difference between the largest and the smallest $\log_{10} E$ over the 5 seeds of one fit, in decades; Range above 1: fits whose seeds differ by more than a factor of ten. Selected best (worst): fits in which the seed with the lowest training loss, the one reported, had the best (worst) held-out error of the five.

Truth	Comparator X	R	95 % interval	Holm p	Failed P/X	Verdict
KS	V-pen	0.827	0.74–0.978	0.024	0/0	differ, size uncertain
	A-hard	0.278	0.155–0.42	0.004	0/0	P better
	F-none	0.0259	0.0102–0.0785	0.004	0/5	P better
DA	V-pen	0.982	0.766–1.21	1.000	0/0	equivalent
	A-hard	114	72.4–181	0.004	0/0	A-hard better
	F-none	0.0561	0.0187–0.163	0.004	0/2	P better
FG	V-pen	1.27	0.266–1.7	1.000	0/3	inconclusive
	A-hard	0.41	0.363–0.545	0.004	0/0	P better
	F-none	0.0322	0.00624–0.122	0.004	0/6	P better
HS	V-pen	0.658	0.528–0.837	0.017	0/0	P better
	A-hard	0.22	0.0913–0.423	0.004	0/0	P better
	F-none	0.0269	0.00575–0.0843	0.004	0/5	P better
TC	V-pen	0.911	0.67–1.17	1.000	0/0	inconclusive
	A-hard	0.023	0.0124–0.0436	0.004	0/0	P better
	F-none	0.00687	0.00198–0.026	0.004	0/4	P better

Table S5. The fifteen isosteric-heat tests on the synthetic truths (Note S2). R: Hodges-Lehmann ratio of the root-mean-square error of Q_{st} at matched loading, P over X, over the 13 materials, with its exact 95 % interval; Holm p over the fifteen tests; Failed: materials on which a task failed (an undefined or non-finite heat, too few states in the loading window), ranked worst.

Dataset	n	H1b V-pen	H1c A-hard	H1d F-none	H1e C1
confirmatory set	13	1.10 [0.93, 1.65] inc	1.51 [0.84, 2.87] inc	0.63 [0.40, 1.22] inc	1.23 [0.92, 1.96] inc
coverage rule, first version	6	0.97 [0.83, 1.27] eq	0.88 [0.32, 1.48] inc	0.46 [0.32, 1.61] inc	1.11 [0.74, 3.20] inc
coverage rule, second version	9	1.19 [0.94, 3.32] inc	1.76 [0.73, 3.92] inc	1.00 [0.49, 2.48] inc	1.49 [0.95, 2.95] inc
no coverage rule	23	1.05 [0.97, 1.20] eq	1.52 [1.06, 2.19] inc	0.59 [0.42, 0.79] P	1.16 [0.97, 1.58] inc
7 points per isotherm	14	1.05 [0.89, 1.54] inc	1.29 [0.73, 2.67] inc	0.63 [0.41, 1.04] inc	1.20 [0.91, 1.66] inc
Co_2Cl_2(BTDD) from its characteristic curves	13	1.04 [0.93, 1.58] inc	1.51 [0.84, 2.87] inc	0.69 [0.40, 1.32] inc	1.28 [0.96, 2.03] inc
digitisation shift, copy 1	13	1.02 [0.90, 1.77] inc	1.35 [0.86, 3.06] inc	0.58 [0.38, 1.07] inc	1.35 [0.94, 2.22] inc
digitisation shift, copy 2	13	1.15 [0.96, 1.72] inc	1.34 [0.72, 2.29] inc	0.53 [0.37, 0.78] P	1.35 [0.97, 1.82] inc

Table S6. The H1 comparisons on the pre-registered dataset variants: R with its exact 95 % interval and the verdict (P, P better; X, comparator better; eq, equivalent within 25 %; inc, inconclusive). The coverage variants are earlier versions of the coverage rule and the dataset without it; the digitisation shifts move every digitised isotherm by ± 0.002 in p/p_{sat}.

Comparison	X	R	Bootstrap 95 %	R, MOFs ($n = 5$)	R, COFs ($n = 8$)
H1b	V-pen	1.10	0.94–1.58	1.09	1.10
H1c	A-hard	1.51	0.87–2.67	0.69	2.67
H1d	F-none	0.63	0.42–1.07	0.43	0.93
H1e	C1	1.23	0.94–1.91	0.99	1.47
H1a (descriptive)	F-hard	0.49	0.33–0.82	0.48	0.50
H1b by source paper	V-pen	1.12	exact 0.92–3.08	10 sources	

Table S7. The H1 comparisons with the percentile bootstrap interval (pre-registered as a secondary), split into MOFs and COFs (pre-registered subgroup, Hodges-Lehmann estimates only, no intervals at these sizes), and, as an exploratory check of the dependence between materials from one paper, H1b with the per-material differences averaged within each source paper (ledger X-4).

Note S1: Can each verdict fail? Synthetic worlds and power

Before the freeze, the whole pipeline (the tuned configurations, fitting from 5 seeds, the violation counts and the analysis code) was run in synthetic worlds that share no uptake data with the confirmatory set. Each world has 13 materials; material i uses the temperatures of confirmatory material i. An expected outcome was written for each world before it ran. In KS0 the truth is a two-site Sips isotherm with van't Hoff affinities, inside the family of P; in DA it is a Dubinin-Astakhov curve, inside the family of the hard characteristic curve (A-hard); in FW it is a Kirchhoff-Sips isotherm with a large heat capacity, so that van't Hoff is wrong; FLAT has one weakly curved Langmuir site; SUB has a sub-linear Sips isotherm outside the family of P and inside that of the penalty model V-pen.

The comparison of P with A-hard discriminated in both directions: P extrapolated better in KS0 ($R = 0.52$, exact 95 % interval 0.35 to 0.70) and A-hard better in DA ($R = 1.26$, 1.07 to 1.49). Three expectations were missed. In FW the comparison of P with the free temperature network F-hard still read "P better" ($R = 0.33$): with two training temperatures the free network cannot extrapolate, so this comparison cannot fail and is reported descriptively. In FLAT the unconstrained network F-none violated physics "often". In SUB, built for V-pen to win, P was slightly better ($R = 0.96$, 0.94 to 0.99); the mixture of Langmuir terms in P approximates a sub-linear isotherm over the measured range. The accuracy comparison of hard constraints with penalties therefore could not be made to fail synthetically, and it is tested as non-inferiority. The rival claim, that penalties suffice for accuracy, predicts the same equivalence, so accuracy alone cannot separate the two approaches; admissibility (H3) is the comparison that can.

The power of each verdict was simulated with the frozen analysis code at $n = 13$ (2000 simulations per setting). At the spread expected for the comparison with V-pen (0.158 decades), the non-inferiority criterion is met with probability 0.0225 when the true ratio is at the margin, $R = 1.25$, and 0.51 (normal) or 0.74 (Student t, 3 degrees of freedom) when the two models are equally accurate. A "not met" outcome is therefore likely even under equal accuracy and is never read as evidence that the constraints cost accuracy.

Note S2: Synthetic truths for the isosteric heat

The heat test uses five truth families with exact heats. Their shape parameters are fitted by bounded least squares to each confirmatory material's middle-temperature isotherm; their thermodynamic parameters are fixed in advance: KS, a Langmuir site with $Q_0 = 55$ kJ mol^{-1} plus a cooperative Sips site with $Q_0 = 48$ kJ mol^{-1}, a heat capacity of 40 J mol^{-1} K^{-1} and an exponent of at least 1.5, inside the family of the Kirchhoff variant of P and near that of P; DA, a Dubinin-Astakhov curve $q = W_0 \exp[-(A/E)^n]$ with $1 \le n \le 6$, inside the family of the characteristic curve; FG, a Fowler-Guggenheim isotherm with $Q_0 = 45$ kJ mol^{-1} and a lateral attraction between 0 and $3.9\,RT_{\min}$, inside neither; HS, strong hydrophilic sites, a Langmuir site with $Q = 80$ kJ mol^{-1} holding at most 0.3 q_{ref} plus a Sips site with $Q = 46$ kJ mol^{-1}, inside the family of P with a heat that falls steeply with loading; TC, a Sips isotherm with $Q = 50$ kJ mol^{-1} whose capacity falls by 0.2 % per kelvin, inside the family of P and outside that of the characteristic curve. Gaussian noise of 0.4 % of q_{ref} is added with a fixed seed per material and family.

Note S3: Digitisation error

The digitiser was tested against published figures whose numeric data are also published, before any model was fitted (logs/digitisation_error.json, logs/digitisation_validation.json). On the vector figure of Mg-SU-102[1] the extraction from the PDF drawing commands recovered 198 of 198 points with a root-mean-square uptake error of 7×10^{-6} g g^{-1} and a largest error of

5×10^{-5} g g^{-1}. The same figure rasterised at 150, 300 and 600 dpi and read by template matching gave root-mean-square uptake errors of 6.3×10^{-4}, 4.0×10^{-4} and 3.5×10^{-4} g g^{-1}, with no false detection; markers drawn hollow are not read by the raster method. On raster figures of HKUST-1 from another paper[2], the root-mean-square uptake error per series was 0.11 to 0.19 % of the axis span, or 0.0012 to 0.0034 of the series' largest uptake, again with no false detection; markers hidden under other markers are missed (recall of the visible markers 0.91 to 1.00). On vector figures of two further papers[3, 4] read through their legend colours, every point was recovered and the root-mean-square uptake error was 0.04 to 0.08 % of the axis span. The noise level of the synthetic tests, 0.4 % of q_{ref}, is the largest raster error per isotherm maximum rounded up.

Note S4: Model definitions

All models map $s = \ln(p/\mathrm{Pa})$ and T to the uptake q (mol kg^{-1}); the code is src/p3a/models.py, and the tuned settings and parameter counts are in Table S8. $T_{\mathrm{ref}} = 300$ K throughout. The same symbol q is used for the uptake of every model; Figure 1a of the main text writes q^* for P's uptake as a derivative of ψ.

P (PotentialIsotherm, van't Hoff). The reduced spreading pressure is

$$\psi(s,T) = \sum_k \frac{c_k(T)}{g_k}\,\mathrm{softplus}\big(g_k(s-\theta_k(T))\big), \qquad q = \frac{\partial\psi}{\partial s} = \sum_k c_k(T)\,\sigma\big(g_k(s-\theta_k(T))\big),$$

with σ the logistic function. Each term is therefore a Langmuir term ($g_k = 1$, the n_{H} Henry terms) or a cooperative Sips-type term ($g_k = 1.5 + \mathrm{softplus}(\cdot)$, the n_{c} cooperative terms) in $K_k(T)p$, with $\ln K_k = -\theta_k$. The capacities are $c_k(T) = c_{k0}\exp[-\alpha_k(T - T_{\mathrm{ref}})]$ with $c_{k0} > 0$ (a softplus), and the positions follow van't Hoff, $\theta_k(T) = \theta_{k0} - (Q_k/R)(1/T - 1/T_{\mathrm{ref}})$. The heats Q_k and the expansion coefficients α_k are unconstrained in sign. Positive capacities make q strictly increasing in s; the Henry terms give $q/p \to \sum_{\mathrm{Henry}} c_k K_k > 0$ as $p \to 0$; the heat $Q_{\mathrm{st}} = -RT^2(\partial q/\partial T)/(\partial q/\partial s)$ has a positive denominator and is finite everywhere, but its sign and size are not constrained. P is thus a positive mixture of Langmuir and cooperative terms with van't Hoff affinities, written as the derivative of one convex potential; C1 is a two-term member of the same broad family with free exponents.

V-none and V-pen. The same function, the same initial function for a given seed and the same tuned structure as P, but with signed capacities and every exponent free (including those of the Henry terms). V-pen adds the penalties below with weight λ.

P-K. As P, with a constant heat capacity Δc_k per term: $\theta_k(T) = \theta_{k0} - (Q_k/R)(1/T - 1/T_{\mathrm{ref}}) - (\Delta c_k/R)[\ln(T/T_{\mathrm{ref}}) + T_{\mathrm{ref}}/T - 1]$.

F-hard. The same potential and shape constraints as P, but $c_k(T)$ and $\theta_k(T)$ are outputs of a one-hidden-layer tanh network of the scaled inverse temperature $\hat{\tau} = (T_{\mathrm{ref}}/T - 1)/0.1$, with each θ_k kept within a bounded range around the data.

A-hard, A-none, A-pen. The characteristic curve $q = W(A)$, $A = RT\ln(p_{\mathrm{sat}}(T)/p)$, with $W(A) = \sum_k c_k\,\sigma(-g_k(A/A_{\mathrm{s}} - A_k))$, $A_{\mathrm{s}} = 2$ kJ mol^{-1}. In A-hard, $c_k > 0$ and $g_k > 0$, so W decreases with A and q increases with p at every temperature; the model is temperature-invariant in A by construction, has no Henry limit, and its heat is exactly $\Delta_{\mathrm{vap}}H(T) + A$. A-none has signed c_k and free g_k; A-pen adds the monotonicity penalty.

F-none, F-pen, F-rel. A fully connected tanh network with inputs $((s - s_c)/s_w, \hat{\tau})$ and one linear output scaled by the largest training uptake (F-none, F-pen), or with $\ln(p/p_{sat})$ in place of s (F-rel); width and depth in Table S8.

C1. $q = \sum_{i=1}^{2} q_{m,i}\,\sigma(g_i(s + \ln b_i(T)))$, that is $q_{m,i}(b_i p)^{g_i}/[1 + (b_i p)^{g_i}]$, with $q_{m,i} > 0$, $g_i = 0.2 + \mathrm{softplus}(\cdot)$ and $\ln b_i(T) = \ln b_{i,\mathrm{ref}} + (Q_i/R)(1/T - 1/T_{\mathrm{ref}})$.

P-sign (exploratory; ledger X-6). P with $Q_k = \mathrm{softplus}(\cdot) > 0$ and $\alpha_k = 10^{-3}\,\mathrm{softplus}(\cdot) \geq 0\ \mathrm{K}^{-1}$, everything else unchanged; the constructor draws the same random numbers as P, so a seed gives the same initial function (checked: maximum difference zero). With both signs fixed, every term's affinity and capacity decrease with temperature, so the heat is positive by construction. The variant was designed after the confirmatory results and fitted on all three tasks (scripts/explore_p_signs.py).

Initial values and data handling. The heats start at $Q_k = 50\ \mathrm{kJ\,mol^{-1}}$ plus a seeded jitter of standard deviation $5\ \mathrm{kJ\,mol^{-1}}$, the expansion coefficients at zero, and the term positions at evenly spaced offsets around the centre of the training pressures. Saturation pressures are from the IAPWS equation[5]. Uptakes were converted to $\mathrm{mol\,kg^{-1}}$ with $M = 18.01528\ \mathrm{g\,mol^{-1}}$ ($\mathrm{g\,g^{-1}}$, wt %) and $22.414\ \mathrm{L\,mol^{-1}}$ ($\mathrm{cm^3(STP)\,g^{-1}}$). In I-mid with an even number of temperatures, the lower of the two middle temperatures is held out.

Penalties. On a collocation grid of 60 pressures (1 Pa to $0.9\,p_{sat}$) at 9 temperatures from $T_{min} - 15$ K to $T_{max} + 15$ K, the monotonicity penalty is the mean of $\max(0, -\partial q/\partial s)^2$ divided by the squared uptake scale. The Henry penalty, at 0.01 Pa, is the mean of $[(\partial q/\partial s - q)/(|\partial q/\partial s| + |q|)]^2$ plus the same relative measure of negative uptake; it is zero only for an isotherm linear in p. The characteristic-curve models have no Henry penalty.

Fitting. The loss is the mean squared error of q divided by the largest training uptake, plus λ times the penalty and an optional ℓ_2 term; Adam (learning rate 10^{-2}) for the number of steps in Table S8, then L-BFGS for 300 iterations, in double precision from 5 seeds, keeping the seed with the lowest training loss. Inputs are centred on the training mean of s and scaled by its standard deviation (at least 0.5).

Model	Role	Form	Tuned settings	Parameters
P	confirmatory reference	potential, van't Hoff	n_{H} 2, n_{c} 2, Adam steps 2000	18
V-pen	H1b (primary)	as P, constraints as penalty	n_{H} 2, n_{c} 2, λ 0.001, Adam steps 2000	20
V-none	descriptive	as P, unconstrained	n_{H} 2, n_{c} 2, Adam steps 2000	20
A-hard	H1c	characteristic curve	n 6, Adam steps 4000	18
A-pen	H3 only	as A-hard, penalty	n 6, λ 0.001, Adam steps 4000	18
A-none	H3 only	as A-hard, unconstrained	n 6, Adam steps 4000	18
F-none	H1d	network in $(s, 1/T)$	width 32, depth 3, ℓ_2 0.0001, Adam steps 4000	2241
F-pen	H3 only	as F-none, penalty	width 32, depth 3, ℓ_2 0.0001, λ 0.1, Adam steps 4000	2241
F-hard	H1a (descriptive)	potential, network in $1/T$	n_{H} 1, n_{c} 2, width 8, Adam steps 2000	72
C1	H1e	dual-site Sips, van't Hoff	Adam steps 4000	8
F-rel	descriptive	network in $(\ln p/p_{\mathrm{sat}}, 1/T)$	width 32, depth 3, Adam steps 2000	2241
P-K	descriptive	P with heat capacity	n_{H} 2, n_{c} 2, Adam steps 2000	22

Table S8. The models, their role in the analysis, the settings chosen on the development set (Methods), and the number of fitted parameters. n_{H}, n_{c}: Henry and cooperative terms; n: terms of the characteristic curve; width and depth of a network; ℓ_2: weight decay; λ: penalty weight. Every fit then ran L-BFGS for 300 iterations.

References

[1] Oppenheim, J.; Yang, Z.; Dinakar, B.; Dincă, M. High-capacity water sorbent cycles without hysteresis under dry conditions. *Nat. Commun.* **2025**, *16* (1), 4297. DOI: 10.1038/s41467-025-59551-2.

[2] Kang, M. S.; Heo, I.; Park, S. H.; Bae, J.; Kim, S.; Kim, G.; Kim, B.-H.; Jeong, N. C.; Yoo, W. C. Time-efficient atmospheric water harvesting using fluorophenyl oligomer incorporated MOFs. *Nat. Commun.* **2024**, *15* (1), 9793. DOI: 10.1038/s41467-024-53853-7.

[3] Liu, Z.; Xu, J.; Xu, M.; Huang, C.; Wang, R.; Li, T.; Huai, X. Ultralow-temperature-driven water-based sorption refrigeration enabled by low-cost zeolite-like porous aluminophosphate. *Nat. Commun.* **2022**, *13* (1), 193. DOI: 10.1038/s41467-021-27883-4.

[4] Gao, M.-Y.; Bezrukov, A. A.; Eaby, A. C.; Deng, C.; Liu, L.; Song, B.-Q.; Nikkhah, S. J.; Vandichel, M.; Zaworotko, M. J. Molecular-level insight into water adsorption and projected atmospheric water harvesting performance in a hydrolytically stable MOF. *Nat. Commun.* **2026**, *17* (1), 9494. DOI: 10.1038/s41467-026-76301-0.

[5] Wagner, W.; Pruss, A. International equations for the saturation properties of ordinary water substance. Revised according to the International Temperature Scale of 1990. *J. Phys. Chem. Ref. Data* **1993**, *22* (3), 783–787. DOI: 10.1063/1.555926.